

HVMX Project Plan

Resumo Executivo

Total LOC: ~5,000

Duração: 3 sprints (8-12 semanas)

Status: Ready to start

Cronograma de Sprints

Sprint 1: Core + MVP (4 semanas)

Objetivo: Runtime básico funcionando com Vulkan

Arquivo	LOC	Status	Testes
crates/hvmx-core/src/lib.rs	200	 Spec	 5 tests
crates/hvmx-core/src/port.rs	150	 Spec	 8 tests
crates/hvmx-core/src/net.rs	250	 Spec	 10 tests
crates/hvmx-core/src/interact.rs	300	 Spec	 15 tests
crates/hvmx-core/src/numb.rs	150	 Spec	 6 tests
crates/hvmx-core/src/book.rs	150	 Spec	 4 tests
crates/hvmx-jit/src/lib.rs	300	 Spec	 8 tests
crates/hvmx-jit/src/ir.rs	400	 Spec	 12 tests
crates/hvmx-jit/src/runtime.rs	300	 Spec	 10 tests
crates/hvmx-jit/src/backend/vulkan.rs	600	 Spec	 20 tests
SUBTOTAL Sprint 1	2,800	Ready	98 tests

Deliverables:

-  Core runtime compilando
-  Backend Vulkan funcional
-  Suite de testes básica
-  CI/CD pipeline

- Benchmarks preliminares
-

Sprint 2: Backends + Memory (3 semanas)

Objetivo: Multi-platform support + otimizações de memória

Arquivo	LOC	Status	Testes
crates/hvmx-jit/src/backend/metal.rs	500	Spec	18 tests
crates/hvmx-jit/src/backend/cuda.rs	400	Spec	15 tests
crates/hvmx-jit/src/backend/cpu.rs	200	Spec	8 tests
crates/hvmx-jit/src/optimizer.rs	300	Spec	12 tests
crates/hvmx-memory/src/lib.rs	150	Spec	5 tests
crates/hvmx-memory/src/unified.rs	300	Spec	15 tests
crates/hvmx-memory/src/prefetch.rs	200	Spec	8 tests
crates/hvmx-memory/src/tile.rs	150	Spec	6 tests
SUBTOTAL Sprint 2	2,200	Ready	87 tests

Deliverables:

- Metal backend (iOS/macOS)
 - CUDA backend (NVIDIA)
 - CPU fallback
 - Unified memory allocator
 - Prefetching strategies
 - Cross-platform benchmarks
-

Sprint 3: Scheduler + Production (3 semanas)

Objetivo: Heterogeneous scheduling + hardening

Arquivo	LOC	Status	Testes
crates/hvmx-scheduler/src/lib.rs	150	Spec	5 tests
crates/hvmx-scheduler/src/partition.rs	250	Spec	10 tests
crates/hvmx-scheduler/src/adaptive.rs	200	Spec	8 tests
crates/hvmx-cli/src/main.rs	150	Spec	6 tests
crates/hvmx-cli/src/commands/*.rs	150	Spec	9 tests
crates/hvmx-cli/src/ui.rs	100	Spec	4 tests
examples/*.rs	300	Spec	8 tests
benches/*.rs	400	Spec	12 tests
tests/*.rs	300	Spec	16 tests
SUBTOTAL Sprint 3	2,000	Ready	78 tests

Deliverables:

- Heterogeneous scheduler (CPU+GPU)
- Adaptive routing
- CLI completo
- 4 exemplos funcionais
- Documentação completa
- Release 1.0

Métricas por Sprint

Sprint	LOC Planejado	LOC Acumulado	Testes	Coverage
1	2,800	2,800	98	85%
2	2,200	5,000	185	87%
3	2,000	7,000	263	90%

Estratégia de Testes

Unit Tests (60%)

- Core: Port, Net, Interact operations
- JIT: IR generation, optimization passes
- Memory: Allocation, prefetching, coherency
- Backends: Mock GPU, CPU fallback

Integration Tests (30%)

- End-to-end reduction
- Multi-backend switching
- Memory management
- CLI workflows

Benchmarks (10%)

- Latency: Small graphs (< 1ms)
 - Throughput: Large graphs (> 100ms)
 - Comparison: CPU vs Vulkan vs Metal vs CUDA
 - Energy: Power efficiency metrics
-

Estrutura de Entrega

Cada Sprint Entrega:

```
sprint-N/
├── code/
│   ├── implemented/  # Código pronto
│   ├── tests/        # Testes passando
│   └── docs/         # Documentação atualizada
└── metrics/
    ├── loc_report.md  # LOC por arquivo
    ├── coverage.html  # Cobertura de testes
    └── benchmarks.json # Resultados de performance
└── artifacts/
    └── CHANGELOG.md  # O que mudou
```

```
└── ROADMAP.md    # Próximos passos  
└── release-notes/ # Notas de release
```

🎯 Critérios de Sucesso

Sprint 1 ✅

- Core runtime compila sem warnings
- Backend Vulkan roda em Android/Linux
- Benchmarks mostram speedup > 2x vs CPU
- CI passa em 3 plataformas

Sprint 2 ✅

- 3+ backends funcionando
- Zero-copy em SoCs mobile
- Benchmarks em 5+ dispositivos
- Documentação API completa

Sprint 3 ✅

- Scheduler adaptativo funcional
 - CLI user-friendly
 - 4 demos funcionando
 - Release candidate 1.0
 - Site hvmx.dev no ar
-

🚀 Quick Start (Para Começar Agora)

```
bash
```

```

# 1. Criar estrutura
mkdir -p hvmx/{crates/{hvmx-core,hvmx-jit,hvmx-memory,hvmx-scheduler,hvmx-cli},benches,examples,tests,docs}

# 2. Inicializar workspace
cargo new --lib hvmx
cd hvmx

# 3. Criar Cargo.toml do workspace
cat > Cargo.toml << 'EOF'
[workspace]
members = [
    "crates/hvmx-core",
    "crates/hvmx-jit",
    "crates/hvmx-memory",
    "crates/hvmx-scheduler",
    "crates/hvmx-cli",
]
resolver = "2"

[workspace.package]
version = "1.0.0"
edition = "2021"
license = "MIT OR Apache-2.0"
EOF

# 4. Inicializar crates
for crate in hvmx-core hvmx-jit hvmx-memory hvmx-scheduler hvmx-cli; do
    cargo init --lib crates/$crate
done

# 5. Setup CI
mkdir -p .github/workflows
# (arquivo CI será criado no próximo passo)

# 6. Começar Sprint 1
# (arquivos serão criados incrementalmente)

```



Próximos Passos Immediatos

1. **Criar repositório GitHub** → (<https://github.com/scoobiii/hvmx>)
2. **Implementar Sprint 1 - Semana 1** → (**hvmx-core**) básico
3. **Setup CI/CD** → GitHub Actions

4. Primeiro benchmark → CPU vs Vulkan

Dashboard de Progresso

Componente	LOC	Status	Testes	Coverage
hvmx-core	1,200	 0%	0/48	0%
hvmx-jit	2,000	 0%	0/50	0%
hvmx-memory	800	 0%	0/34	0%
hvmx-scheduler	600	 0%	0/23	0%
hvmx-cli	400	 0%	0/19	0%
TOTAL	5,000	0%	0/174	0%

 = Completo |  = Em progresso |  = Não iniciado

Ready to start Sprint 1? 